

# chess

Chess bot written in rust

## Improvement Roadmap

### Speed Optimizations

#### Incremental PST Evaluation ✓ DONE

Currently the evaluation function iterates all pieces on every leaf node, which is  $O(n\_pieces)$  per eval call. Now maintains running PST scores that update incrementally during make\_move / unmake\_move . Evaluation is now  $O(1)$  per position.

#### Remove pieces: Vec<Piece> Redundancy ✓ PARTIAL

Added iter\_pieces() method using bitboards. Updated hot paths (evaluation, material counting, pin detection) to use bitboards. The pieces Vec is still used in make/unmake\_move but is no longer on the critical evaluation path.

#### Lazy/Incremental Attack Map Updates ! ATTEMPTED

recompute\_attack\_maps() is called on every make\_move and unmake\_move , regenerating all sliding piece attacks from scratch. Lazy evaluation was attempted using atomics for thread-safety but the atomic overhead (~30-40% perf slowdown) outweighed the benefits. Future approaches: incremental updates that only recalculate affected sliding piece rays, or magic bitboards which would make full recomputation fast enough.

#### Stack-Allocated Move Lists

get\_legal\_moves() allocates a new Vec<Move> per node, causing heap allocation pressure. Replace with ArrayVec<Move, 256> (stack-allocated) or a thread-local move list pool. The maximum legal moves in any chess position is 218, so a fixed-size array suffices.

#### Magic Bitboards for Sliding Pieces

Currently using classical ray attacks with blocker detection via trailing\_zeros / leading\_zeros . Magic bitboards use a precomputed lookup table indexed by (occupancy \* magic) >> shift , reducing sliding piece attack generation to a single array lookup. Typically provides ~2x speedup for move generation.

### Strength Optimizations

#### Enable Endgame PST Tables ✓ DONE

The codebase defines `PAWNS_END` and `KING_END` tables but `is_endgame` is hardcoded to `false`. Now uses tapered evaluation with phase detection:

`phase = (queens*4 + rooks*2 + bishops + knights)`. Interpolates between middlegame and endgame PST scores based on remaining material.

**Killer Move Heuristic** ✓ DONE

Track the last 2 quiet moves that caused a beta cutoff at each ply depth. Implemented in `SearchState` struct. Stores 2 killer moves per ply, prioritized in move ordering after TT move and captures (90,000 for first killer, 80,000 for second).

**History Heuristic** ✓ DONE

Maintain a `[Color] [From] [To]` table that accumulates a bonus each time a quiet move causes a beta cutoff, scaled by depth<sup>2</sup>. Implemented in `SearchState`. History scores indexed by `[color] [from_sq] [to_sq]`, used for quiet move ordering. Scores aged at each depth iteration to gradually forget old information.

**Principal Variation Search (PVS)** ✓ DONE

After searching the first move with a full window, search remaining moves with a null window (`-alpha-1, -alpha`). If the null window search fails high, re-search with the full window. Implemented in `negamax_with_tt_mut`. First move uses full window, subsequent moves use null window (`-alpha-1, -alpha`), with full re-search on fail high. Combined with LMR for later quiet moves. Benchmarks show ~12-28% faster search at depths 5-6.

**Aspiration Windows** ✓ DONE

In iterative deepening, use a narrow window centered on the previous iteration's score (e.g.,  $\pm 25$  centipawns) instead of  $(-\infty, +\infty)$ . Implemented in `iterative_deepening` and `iterative_deepening_timed`. After depth 1, uses  $\pm 25$ cp window centered on previous score. On fail-low/fail-high, doubles window and re-searches. Falls back to full window if window exceeds  $\pm 200$ cp.

**Check Extensions** ✓ DONE

When a move gives check, extend the search depth by 1 ply. Implemented in `negamax_with_tt_mut`. After making a move, check if opponent is in check; if so, extend search depth by 1. Also prevents LMR from being applied to moves that give check. This avoids horizon effects where forcing check sequences are cut off prematurely.

**Static Exchange Evaluation (SEE)** ✓ DONE

Before searching a capture in quiescence, simulate the full exchange sequence on that square to determine if the capture is winning, equal, or losing. Implemented `see()` function that simulates exchange sequences. In quiescence search, prunes losing captures when `attacker_value > victim_value`. Uses least-valuable-attacker ordering for accurate exchange simulation.

**Futility Pruning** ✓ DONE

At low depths (1-2 ply from leaf), if the static evaluation plus a margin is still below alpha, skip searching quiet moves entirely. Implemented in `negamax_with_tt_mut`. At depths 1-2, skips

quiet moves (non-captures, non-promotions) if static\_eval + margin < alpha. Margins: depth 1 = 200cp, depth 2 = 500cp. Not applied when in check or for the first move.

## Improved Evaluation Terms

The current eval is material + PST only. Add:

- **Pawn structure:** Penalize doubled pawns (-10), isolated pawns (-20), bonus for passed pawns (+20 to +100 by rank)
- **Mobility:** Count pseudo-legal moves per piece, bonus for more active pieces
- **King safety:** Bonus for pawn shield in front of castled king, penalty for open files near king
- **Bishop pair:** +30 bonus for having both bishops

## Implementation Priority

Phase	Task	Type	Expected Impact	Status
1	Enable endgame PST interpolation	Strength	Quick win, minimal code	✓ Done
2	Incremental PST evaluation	Speed	15-20% faster	✓ Done
3	Remove pieces Vec	Speed	10-15% faster	✓ Partial
4	Killer moves + history heuristic	Strength	Much better move ordering	✓ Done
5	Lazy/incremental attack maps	Speed	10-20% faster	⚠ Attempted
6	PVS search	Strength	Better node efficiency	✓ Done
7	Check extensions	Strength	Avoids horizon effects	✓ Done
8	Magic bitboards	Speed	~2x faster movegen	
9	SEE pruning	Strength	Smaller quiescence tree	✓ Done
10	Futility pruning	Strength	Smaller search tree	✓ Done
11	Improved eval terms	Strength	Better positional play	